

## 03\_asl\_cnn

August 4, 2024

### 1 Python Program to Implement a CNN with TensorFlow and Keras

```
[ ]: import tensorflow as tf
from tensorflow.keras import layers, models

# Assuming X_train, y_train, X_valid, y_valid are already defined

# Define the CNN model
model = models.Sequential()

# First Convolutional Layer
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(64, 64, 1)))
model.add(layers.BatchNormalization())
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Dropout(0.25))

# Second Convolutional Layer
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.BatchNormalization())
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Dropout(0.25))

# Flatten Layer
model.add(layers.Flatten())

# Dense Layer
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dropout(0.5))

# Output Dense Layer
model.add(layers.Dense(10, activation='softmax'))

# Compile the model
model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
```

```

        metrics=['accuracy']
    )

    # Fit the model
    history = model.fit(
        X_train, y_train,
        epochs=10,
        batch_size=32,
        validation_data=(X_valid, y_valid)
    )

    # Print the model summary
    print(model.summary())

```

## 2 Convolutional Neural Networks

In the previous lecture, we built and trained a simple model to classify ASL images. The model was able to learn how to correctly classify the training dataset with very high accuracy, but, it did not perform nearly as well on validation dataset. This behavior of not generalizing well to non-training data is called [overfitting](#), and in this section, we will introduce a popular kind of model called a [convolutional neural network](#) that is especially good for reading images and classifying them.

### 2.1 Loading and Preparing the Data

The below cell contains the data preprocessing techniques we learned in the previous lectures. Execute it.

```

[ ]: import tensorflow.keras as keras
import pandas as pd

train_df = pd.read_csv("sign_mnist_train.csv")
valid_df = pd.read_csv("sign_mnist_valid.csv")

# Separate out our target values
y_train = train_df['label']
y_valid = valid_df['label']
del train_df['label']
del valid_df['label']

# Separate out our image vectors
x_train = train_df.values
x_valid = valid_df.values

# Turn our scalar targets into binary categories
num_classes = 24
y_train = keras.utils.to_categorical(y_train, num_classes)
y_valid = keras.utils.to_categorical(y_valid, num_classes)

```

```
# Normalize our image data
x_train = x_train / 255
x_valid = x_valid / 255
```

## 2.2 Reshaping Images for a CNN

The individual pictures in our dataset are in the format of long lists of 784 pixels:

```
[ ]: x_train.shape, x_valid.shape
```

```
[ ]: ((27455, 784), (7172, 784))
```

In this format, we don't have all the information about which pixels are near each other. Because of this, we can't apply convolutions that will detect features. Let's reshape our dataset so that they are in a 28x28 pixel format. This will allow our convolutions to associate groups of pixels and detect important features.

Note that for the first convolutional layer of our model, we need to have not only the height and width of the image, but also the number of **color channels**. Our images are grayscale, so we'll just have 1 channel.

That means that we need to convert the current shape (27455, 784) to (27455, 28, 28, 1). As a convenience, we can pass the **reshape** method a -1 for any dimension we wish to remain the same, therefore:

```
[ ]: x_train = x_train.reshape(-1,28,28,1)
     x_valid = x_valid.reshape(-1,28,28,1)
```

```
[ ]: x_train.shape
```

```
[ ]: (27455, 28, 28, 1)
```

```
[ ]: x_valid.shape
```

```
[ ]: (7172, 28, 28, 1)
```

```
[ ]: x_train.shape, x_valid.shape
```

```
[ ]: ((27455, 28, 28, 1), (7172, 28, 28, 1))
```

## 2.3 Creating a Convolutional Model

These days, many data scientists start their projects by borrowing model properties from a similar project. Assuming the problem is not totally unique, there's a great chance that people have created models that will perform well which are posted in online repositories like [TensorFlow Hub](#) and the [NGC Catalog](#).

```
[ ]: from tensorflow.keras.models import Sequential
     from tensorflow.keras.layers import (
```

```

Dense,
Conv2D,
MaxPool2D,
Flatten,
Dropout,
BatchNormalization,
)

model = Sequential()
model.add(Conv2D(75, (3, 3), strides=1, padding="same", activation="relu",
                input_shape=(28, 28, 1)))
model.add(BatchNormalization())
model.add(MaxPool2D((2, 2), strides=2, padding="same"))
model.add(Conv2D(50, (3, 3), strides=1, padding="same", activation="relu"))
model.add(Dropout(0.2))
model.add(BatchNormalization())
model.add(MaxPool2D((2, 2), strides=2, padding="same"))
model.add(Conv2D(25, (3, 3), strides=1, padding="same", activation="relu"))
model.add(BatchNormalization())
model.add(MaxPool2D((2, 2), strides=2, padding="same"))
model.add(Flatten())
model.add(Dense(units=512, activation="relu"))
model.add(Dropout(0.3))
model.add(Dense(units=num_classes, activation="softmax"))

```

### 2.3.1 Conv2D

These are our 2D convolutional layers. Small kernels will go over the input image and detect features that are important for classification. Earlier convolutions in the model will detect simple features such as lines. Later convolutions will detect more complex features. Let's look at our first Conv2D layer:

```
model.add(Conv2D(75 , (3,3) , strides = 1 , padding = 'same'...)
```

75 refers to the number of filters that will be learned. (3,3) refers to the size of those filters. Strides refer to the step size that the filter will take as it passes over the image. Padding refers to whether the output image that's created from the filter will match the size of the input image.

### 2.3.2 BatchNormalization

Like normalizing our inputs, batch normalization scales the values in the hidden layers to improve training. [Read more about it in detail here.](#)

### 2.3.3 MaxPool2D

Max pooling takes an image and essentially shrinks it to a lower resolution. It does this to help the model be robust to translation (objects moving side to side), and also makes our model faster.

### 2.3.4 Dropout

Dropout is a technique for preventing overfitting. Dropout randomly selects a subset of neurons and turns them off, so that they do not participate in forward or backward propagation in that particular pass. This helps to make sure that the network is robust and redundant, and does not rely on any one area to come up with answers.

### 2.3.5 Flatten

Flatten takes the output of one layer which is multidimensional, and flattens it into a one-dimensional array. The output is called a feature vector and will be connected to the final classification layer.

### 2.3.6 Dense

We have seen dense layers before in our earlier models. Our first dense layer (512 units) takes the feature vector as input and learns which features will contribute to a particular classification. The second dense layer (24 units) is the final classification layer that outputs our prediction.

## 2.4 Summarizing the Model

This may feel like a lot of information, but don't worry. It's not critical that to understand everything right now in order to effectively train convolutional models. Most importantly we know that they can help with extracting useful information from images, and can be used in classification tasks.

Here, we summarize the model we just created. Notice how it has fewer trainable parameters than the model in the previous notebook:

```
[ ]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 75)	750
batch_normalization (Batch Normalization)	(None, 28, 28, 75)	300
max_pooling2d (MaxPooling2D)	(None, 14, 14, 75)	0
conv2d_1 (Conv2D)	(None, 14, 14, 50)	33800
dropout (Dropout)	(None, 14, 14, 50)	0
batch_normalization_1 (Batch Normalization)	(None, 14, 14, 50)	200
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 50)	0
conv2d_2 (Conv2D)	(None, 7, 7, 25)	11275

```

-----
batch_normalization_2 (Batch Normalization) (None, 7, 7, 25) 100
-----
max_pooling2d_2 (MaxPooling2D) (None, 4, 4, 25) 0
-----
flatten (Flatten) (None, 400) 0
-----
dense (Dense) (None, 512) 205312
-----
dropout_1 (Dropout) (None, 512) 0
-----
dense_1 (Dense) (None, 24) 12312
=====
Total params: 264,049
Trainable params: 263,749
Non-trainable params: 300
-----

```

## 2.5 Compiling the Model

We'll compile the model just like before:

```
[ ]: model.compile(loss="categorical_crossentropy", metrics=["accuracy"])
```

## 2.6 Training the Model

Despite the very different model architecture, the training looks exactly the same. Run the cell below to train for 20 epochs and let's see if the accuracy improves:

```
[ ]: model.fit(x_train, y_train, epochs=10, verbose=1, validation_data=(x_valid,
    ↪ y_valid))
```

```

Epoch 1/10
858/858 [=====] - 5s 5ms/step - loss: 0.3072 -
accuracy: 0.9066 - val_loss: 0.2882 - val_accuracy: 0.9127
Epoch 2/10
858/858 [=====] - 4s 5ms/step - loss: 0.0195 -
accuracy: 0.9937 - val_loss: 0.3995 - val_accuracy: 0.9095
Epoch 3/10
858/858 [=====] - 4s 5ms/step - loss: 0.0102 -
accuracy: 0.9965 - val_loss: 0.3287 - val_accuracy: 0.9223
Epoch 4/10
858/858 [=====] - 4s 5ms/step - loss: 0.0077 -
accuracy: 0.9974 - val_loss: 0.2268 - val_accuracy: 0.9516
Epoch 5/10
858/858 [=====] - 4s 5ms/step - loss: 0.0076 -
accuracy: 0.9982 - val_loss: 0.3375 - val_accuracy: 0.9165
Epoch 6/10
858/858 [=====] - 4s 5ms/step - loss: 0.0069 -

```

```
accuracy: 0.9983 - val_loss: 0.2774 - val_accuracy: 0.9495
Epoch 7/10
858/858 [=====] - 4s 5ms/step - loss: 0.0028 -
accuracy: 0.9992 - val_loss: 0.3765 - val_accuracy: 0.9413
Epoch 8/10
858/858 [=====] - 4s 5ms/step - loss: 0.0037 -
accuracy: 0.9993 - val_loss: 0.1556 - val_accuracy: 0.9787
Epoch 9/10
858/858 [=====] - 4s 5ms/step - loss: 0.0029 -
accuracy: 0.9993 - val_loss: 0.3367 - val_accuracy: 0.9561
Epoch 10/10
858/858 [=====] - 4s 5ms/step - loss: 0.0025 -
accuracy: 0.9994 - val_loss: 0.2632 - val_accuracy: 0.9501
```


```
[ ]: <tensorflow.python.keras.callbacks.History at 0x7fb3c326ab38>
```

## 2.7 Discussion of Results

It looks like this model is significantly improved! The training accuracy is very high, and the validation accuracy has improved as well. This is a great result, as all we had to do was swap in a new model.

## cnn

Convolution neural network(convenet) are a special kind of neural network for processing data that has a known grid-like topology like time series data (1D) or images (2d)




Feature	Artificial Neural Networks (ANNs)	Convolutional Neural Networks (CNNs)
<b>Architecture</b>	Fully connected layers	Convolutional layers followed by pooling layers, then fully connected layers
<b>Connection Pattern</b>	Each neuron connected to every neuron in subsequent layer	Local connections using filters/kernels
<b>Handling of Input Data</b>	Designed for 1D structured data (e.g., tabular data)	Designed for 2D/3D spatial data (e.g., images, videos)
<b>Parameter Sharing</b>	No weight sharing, each connection has its own weight	Weights are shared within the filters
<b>Feature Extraction</b>	Does not inherently extract spatial features	Extracts spatial hierarchies (edges, textures, objects)
<b>Dimensionality Reduction</b>	Not inherently designed for dimensionality reduction	Uses pooling layers to reduce spatial dimensions
<b>Computational Efficiency</b>	Less efficient with high-dimensional data	More efficient due to weight sharing and reduced parameters
<b>Training Complexity</b>	Requires more parameters for high-dimensional data	Fewer parameters due to shared weights and pooling
<b>Applications</b>	Classification, regression, time series prediction	Image and video recognition, object detection, image segmentation
<b>Examples of Use Cases</b>	Predicting stock prices, customer churn prediction	Face recognition, medical image analysis, autonomous driving
<b>Sensitivity to Input</b>	Sensitive to the scale of input features	Less sensitive due to normalization layers (e.g., batch normalization)
<b>Overfitting</b>	Higher risk of overfitting with large datasets	Lower risk due to spatial hierarchies and pooling layers


## why not ANN?


1. high computation cost
2. overfitting
3. loss of imp info like spatial arrangement of pixels






## CNN INTUITION



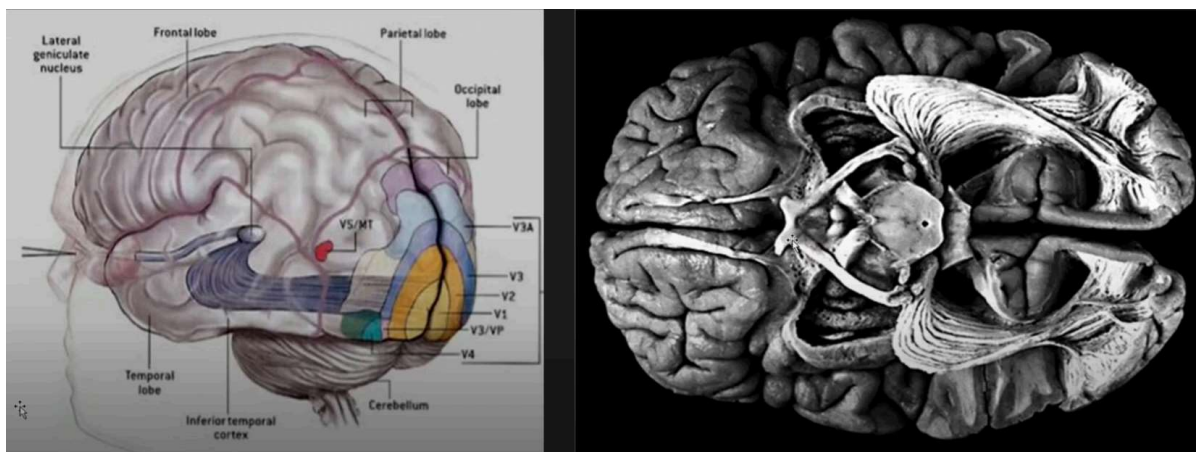


## APPLICATION WISE






## HUMAN VISUAL CORTEX




## EXPRIMENT SETUP






- EDGE DETECTION (PRIMARY FEATURE CELL)
- COMPLEX DETECTION (COMPLEX FEATURE CELL)




## CONVOLUTION OPERATION

- CONV-POOLING-FC




\*\*BASIC OF IMAGES

- GRayscale(B/W) (ONE)
- RGB (THREE)




157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	169	181
206	109	5	124	181	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	105	207	233	233	214	220	239	228	88	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	86	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	96	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	209	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

157	153	174	168	150	152	129	151	172	161	155	156
195	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	169	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	105	207	233	233	214	220	239	228	88	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	86	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	96	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	209	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218



## EDGE DETECTION IN CONVOLUTION



<https://deeplizard.com/resource/pavq7noze2> (<https://deeplizard.com/resource/pavq7noze2>)

- 1 # SHAPE AFTER CONVOLUTION
- 2 ##### IMAGE \* FILTER = FEATURE MAP
- 3 ##### (N\*N) \*(M\*M) =(N-M+1)\*(N-M+1)

$$4 - (28 \times 28) * (3 \times 3) = (26 \times 26)$$


RGB IMAGE \* FILTER = FEATURE MAP (NNC) \* (MMC) = (N-M+1)(N-M+1) (SINGLE CHANNEL (28x28) \* (3x3) = (26x26)

## RGB

IMAGE \* FILTER = FEATURE MAP

(NNC) (MMC) = (N-M+1)(N-M+1) (SINGLE CHANNEL


- (28x28) \* (3x3) = (26x26)



1 # IN OUTPUT LAYER 4\*4\*2 MEANS 2 FILTER APPLY

## PADDING

- padding is a technique used to preserve the spatial dimensions of the input image after convolution operations on a feature map. Padding involves adding extra pixels around the border of the input feature map before convolution.
- For a grayscale (n x n) image and (f x f) filter/kernel, the dimensions of the image resulting from a convolution operation is (n - f + 1) x (n - f + 1).



- p = number of layers of zeros added to the border of the image,
- then (n x n) image → (n + 2p) x (n + 2p) image after padding.
- (n + 2p) x (n + 2p) \* (f x f) → outputs (n + 2p - f + 1) x (n + 2p - f + 1) images

```
In [27]: 1 import tensorflow
2 from tensorflow import keras
3 from keras.layers import Dense, Conv2D, Flatten
4 from keras import Sequential
5 from keras.datasets import mnist
```

```
In [28]: 1 (X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
In [29]: 1 model = Sequential()
2
3 model.add(Conv2D(32, kernel_size=(3, 3), padding='valid', activation='relu',
4 model.add(Conv2D(32, kernel_size=(3, 3), padding='valid', activation='relu'))
5 model.add(Conv2D(32, kernel_size=(3, 3), padding='valid', activation='relu'))
6
7 model.add(Flatten())
8
9 model.add(Dense(128, activation='relu'))
10 model.add(Dense(10, activation='softmax'))
```

```
In [30]: 1 model.summary()
```

Model: "sequential\_8"

Layer (type)	Output Shape	
conv2d_22 (Conv2D)	(None, 26, 26, 32)	
conv2d_23 (Conv2D)	(None, 24, 24, 32)	
conv2d_24 (Conv2D)	(None, 22, 22, 32)	
flatten_8 (Flatten)	(None, 15488)	
dense_16 (Dense)	(None, 128)	
dense_17 (Dense)	(None, 10)	

Total params: 2,002,698 (7.64 MB)

Trainable params: 2,002,698 (7.64 MB)

Non-trainable params: 0 (0.00 B)

```
In [33]: 1 model = Sequential()
2
3 model.add(Conv2D(32,kernel_size=(3,3),padding='same', activation='relu', i
4 model.add(Conv2D(32,kernel_size=(3,3),padding='same', activation='relu'))
5 model.add(Conv2D(32,kernel_size=(3,3),padding='same', activation='relu'))
6
7 model.add(Flatten())
8
9 model.add(Dense(128,activation='relu'))
10 model.add(Dense(10,activation='softmax'))
```

C:\Users\VISHAL\anaconda3\Lib\site-packages\keras\src\layers\convolutional\base\_conv.py:107: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
In [34]: 1 model.summary()
```

Model: "sequential\_10"

Layer (type)	Output Shape	
conv2d_28 (Conv2D)	(None, 28, 28, 32)	
conv2d_29 (Conv2D)	(None, 28, 28, 32)	
conv2d_30 (Conv2D)	(None, 28, 28, 32)	
flatten_10 (Flatten)	(None, 25088)	
dense_20 (Dense)	(None, 128)	
dense_21 (Dense)	(None, 10)	

Total params: 3,231,498 (12.33 MB)

Trainable params: 3,231,498 (12.33 MB)

Non-trainable params: 0 (0.00 B)



```
In [31]: 1 model = Sequential()
2
3 model.add(Conv2D(32,kernel_size=(3,3),padding='same',strides=(2,2), activation='relu'))
4 model.add(Conv2D(32,kernel_size=(3,3),padding='same',strides=(2,2), activation='relu'))
5 model.add(Conv2D(32,kernel_size=(3,3),padding='same',strides=(2,2), activation='relu'))
6
7 model.add(Flatten())
8
9 model.add(Dense(128,activation='relu'))
10 model.add(Dense(10,activation='softmax'))
```

```
In [32]: 1 model.summary()
```

Model: "sequential\_9"


Layer (type)	Output Shape	
conv2d_25 (Conv2D)	(None, 14, 14, 32)	
conv2d_26 (Conv2D)	(None, 7, 7, 32)	
conv2d_27 (Conv2D)	(None, 4, 4, 32)	
flatten_9 (Flatten)	(None, 512)	
dense_18 (Dense)	(None, 128)	
dense_19 (Dense)	(None, 10)	

Total params: 85,770 (335.04 KB)

Trainable params: 85,770 (335.04 KB)

Non-trainable params: 0 (0.00 B)

**sTRIDE:** Stride determines how many squares or pixels our filters skip when they move across the image, from left to right and from top to bottom.



- 1 The input dimensions of the image - > I (ixi)
- 2 The size of filter/kernel - > F (fxf)
- 3 Strides - > S (integer)
- 4 Padding - > P (integer)
- 5 Depth/Number of feature maps/activation maps - > D (integer)
- 6 Convolution Output dimension =  $[(I - F + 2 * P) / S] + 1 \times D$  > Formula1

In [21]:

```

1 model = Sequential()
2
3 model.add(Conv2D(32,kernel_size=(3,3),padding='same',strides=(2,2), activation='relu'))
4 model.add(Conv2D(32,kernel_size=(3,3),padding='same',strides=(2,2), activation='relu'))
5 model.add(Conv2D(32,kernel_size=(3,3),padding='same',strides=(2,2), activation='relu'))
6
7 model.add(Flatten())
8
9 model.add(Dense(128,activation='relu'))
10 model.add(Dense(10,activation='softmax'))

```

```
In [22]: 1 model.summary()
```

Model: "sequential\_6"

Layer (type)	Output Shape	
conv2d_17 (Conv2D)	(None, 14, 14, 32)	
conv2d_18 (Conv2D)	(None, 7, 7, 32)	
conv2d_19 (Conv2D)	(None, 4, 4, 32)	
flatten_6 (Flatten)	(None, 512)	
dense_12 (Dense)	(None, 128)	
dense_13 (Dense)	(None, 10)	



Total params: 85,770 (335.04 KB)

Trainable params: 85,770 (335.04 KB)

Non-trainable params: 0 (0.00 B)

<https://deeplizard.com/resource/pavq7noze3>  
(<https://deeplizard.com/resource/pavq7noze3>)



```
In [23]: 1 import tensorflow
2 from tensorflow import keras
3 from keras.layers import Dense, Conv2D, Flatten, MaxPooling2D
4 from keras import Sequential
5 from keras.datasets import mnist
```

```
In [24]: 1 (X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
In [25]: 1 model = Sequential()
2
3 model.add(Conv2D(32,kernel_size=(3,3),padding='valid', activation='relu',
4 model.add(MaxPooling2D(pool_size=(2, 2), strides=2, padding='valid'))
5 model.add(Conv2D(32,kernel_size=(3,3),padding='valid', activation='relu'))
6 model.add(MaxPooling2D(pool_size=(2, 2), strides=2, padding='valid'))
7
8 model.add(Flatten())
9
10 model.add(Dense(128,activation='relu'))
11 model.add(Dense(10,activation='softmax'))
```

```
In [26]: 1 model.summary()
```

Model: "sequential\_7"

Layer (type)	Output Shape	
conv2d_20 (Conv2D)	(None, 26, 26, 32)	
max_pooling2d_2 (MaxPooling2D)	(None, 13, 13, 32)	
conv2d_21 (Conv2D)	(None, 11, 11, 32)	
max_pooling2d_3 (MaxPooling2D)	(None, 5, 5, 32)	
flatten_7 (Flatten)	(None, 800)	
dense_14 (Dense)	(None, 128)	
dense_15 (Dense)	(None, 10)	



Total params: 113,386 (442.91 KB)

Trainable params: 113,386 (442.91 KB)

Non-trainable params: 0 (0.00 B)